CS186 Discussion 1

(Rendezvous, External Sorting, External Hashing)

Matthew Deng



Matthew Deng

UC Berkeley B.S.'16 EECS matthew.deng@berkeley.edu

Office Hours: Mondays 2-4PM @ Soda 611







Icebreaker

- Get to know the people around you:
 - Name
 - Year
 - Major
 - Interesting fact about yourself

Warm-Up Worksheet #1, 2

1. What is time-space rendezvous and why do we use it?

- 1. What is time-space rendezvous and why do we use it?
 - records are in the same place (RAM) at the same time
 - useful when certain records need to be co-resident but are not guaranteed to be in the same input chunk
 - implemented through out-of-core algorithms
 - external sorting
 - external hashing

2. What is the difference between external sorting and external hashing?

- 2. What is the difference between external sorting and external hashing?
 - Sorting:
 - conquer and merge
 - good if we need to sort or already sorted
 - Hashing:
 - divide and conquer
 - good if we need to remove duplicates
 - Same memory requirement, i/o cost (for 2 passes)

External Sorting

External Sorting

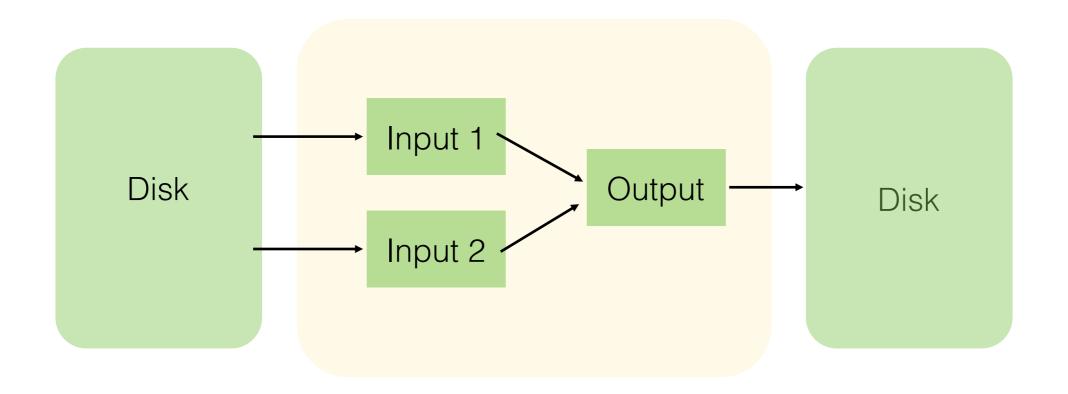
Want to sort data that does not fit in memory

2-Way Merge Sort

- Pass 0 (Conquer):
 - Sort 1 page at a time in memory
- Pass 1, 2, ..., etc. (Merge):
 - Merge 2 runs

2-Way Merge Sort

Credits to Michelle Nguyen



Buffer size of 3 pages

- 3,4
- 6,2
- 9,4
- 8,7
- 5,6
- 6,5
- 1,4
- 4,2

Pass 0

Output

- 3,4
- 6,2
- 9,4
- 8,7
- 5,6
- 6,5
- 1,4
- 4,2

Pass 0

Output

6,2

9,4

8,7

5,6

6,5

1,4

4,2

Pass 0

Output

6.2

0.4

8,7

5,6

6,5

1,4

4,2

2 /

Pass 0

Output

9,4

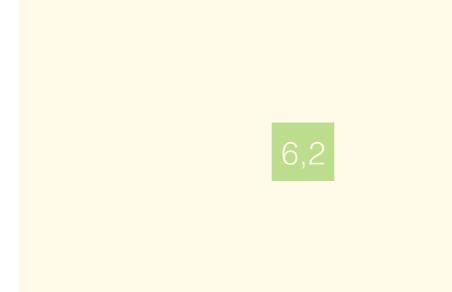
8,7

5,6

6,5

1,4

4,2



Pass 0

Output

9,4

8,7

5,6

6,5

1,4

4,2

2,6

Pass 0

Output

9,4

8,7

5,6

6,5

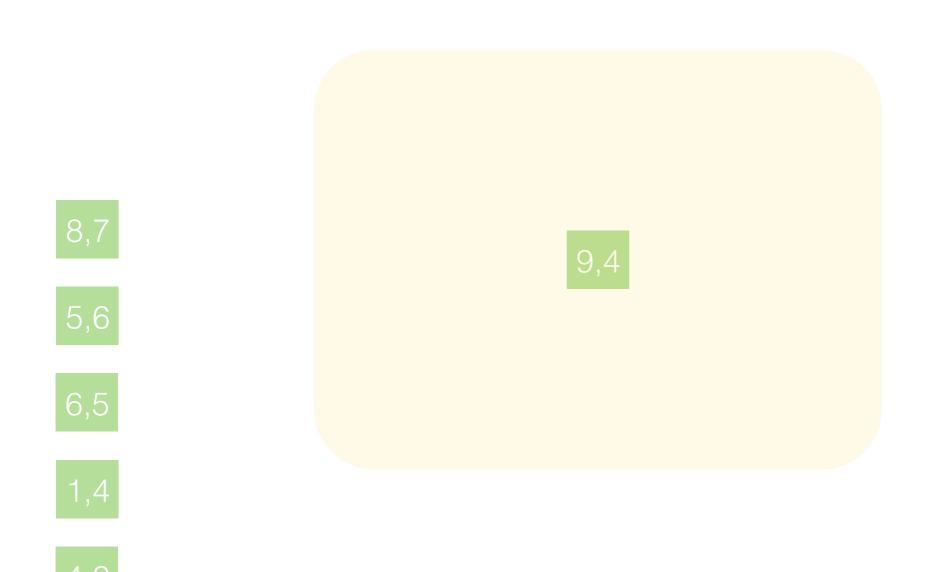
1,4

4,2

3,4

Pass 0

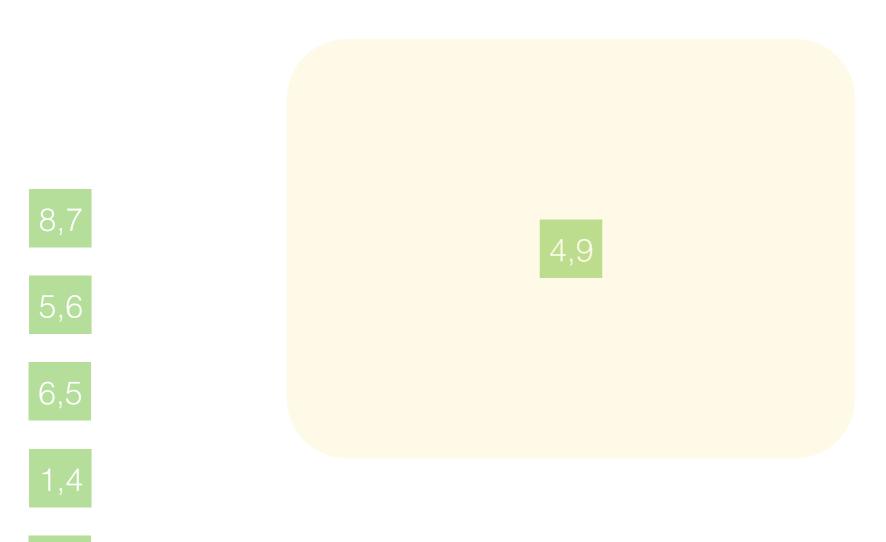
Output



3,4

Pass 0

Output



3,4

2,6

Pass 0

Output

8,7

5,6

6,5

1,4

4,2

3,4

2,6

Pass 0

Output

3,4

2,6

4,9

7,8

5,6

5,6

1,4

Input Pass 0

3,4

3,4

6,2

2,6

9,4

4,9

8,7

7,8

5,6

5,6

6,5

5,6

1,4

1,4

4,2

2,4

1 page runs

Pass 1

Output

3,4

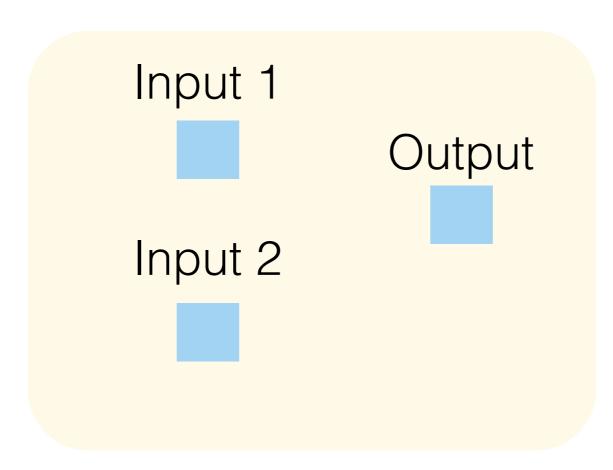
2,6

7,8

5,6

5,6

1,4



Pass 1

Output

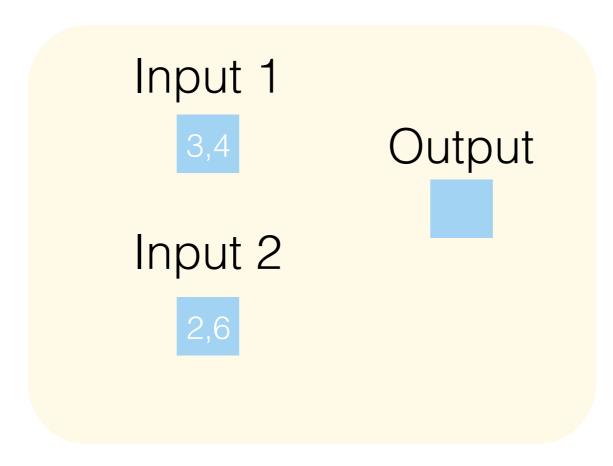
4,9

7,8

5,6

5,6

1,4



Pass 1

Output

4,9

7,8

5,6

5,6

1,4

2,4

Input 1

3,4

Input 2

6

Output

2

Pass 1

Output

4,9

7,8

5,6

5,6

1,4

2,4

Input 1

4

Output

2,3

Input 2

6

Pass 1

Output

2,3

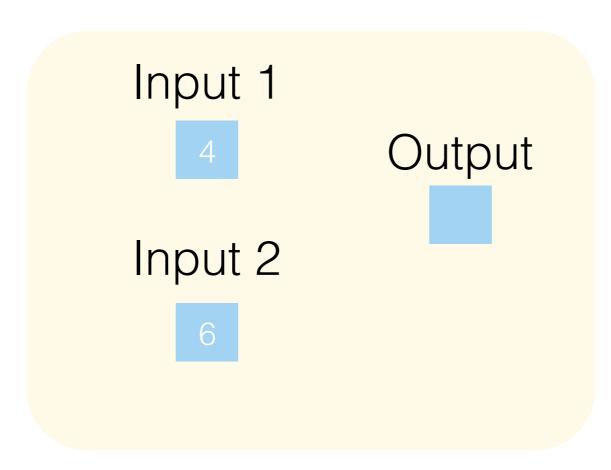
4,9

7,8

5,6

5,6

1,4



Pass 1

Output

2,3

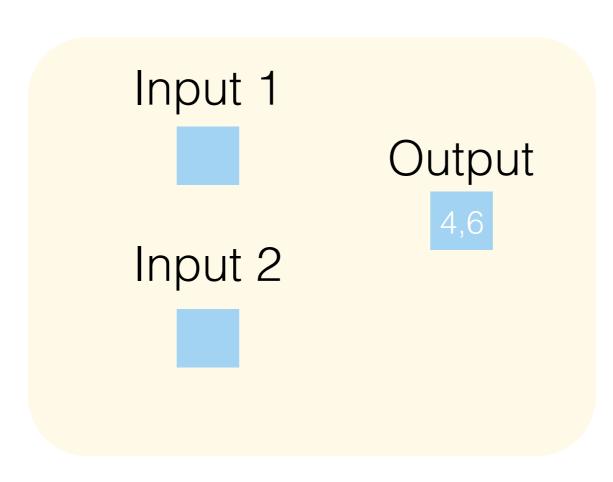
4,9

7,8

5,6

5,6

1,4



Pass 1

Output

2,3

4,6

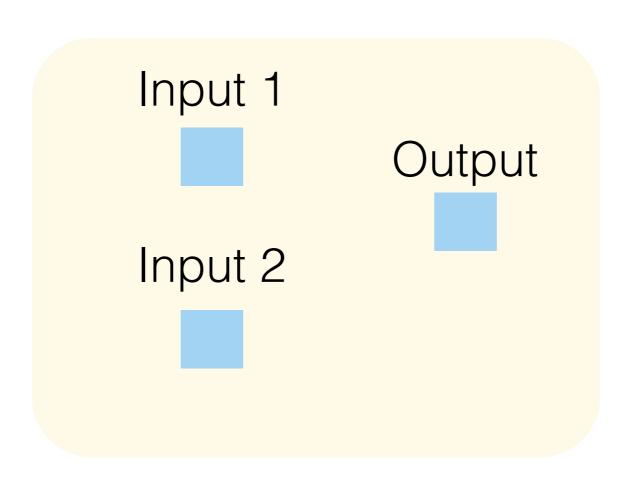
4,9

7,8

5,6

5,6

1,4



Pass 1

Output

2,3

4,6

Input 1

4,9

Output

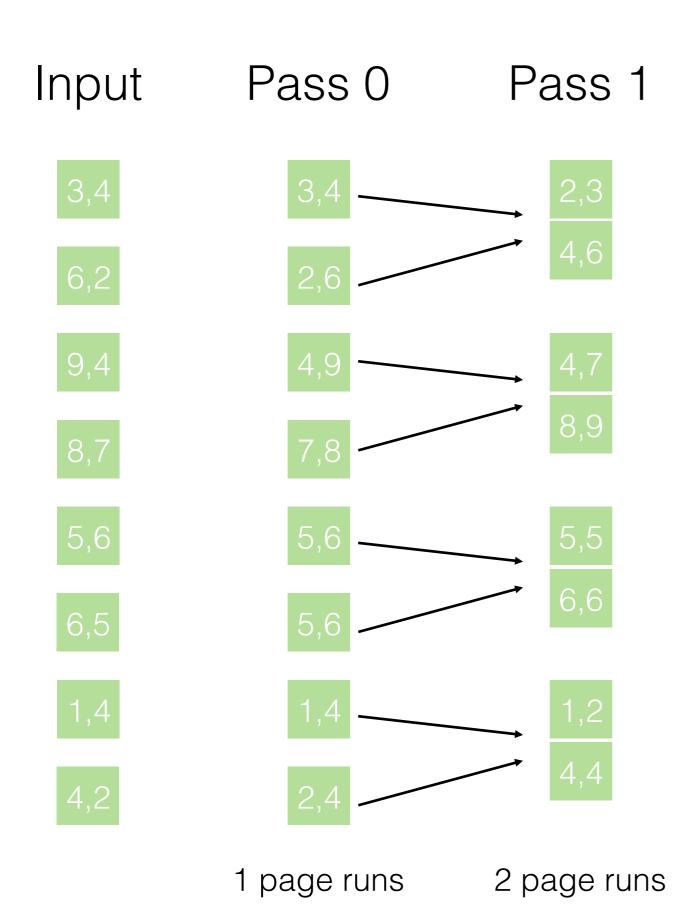
Input 2

7,8

5,6

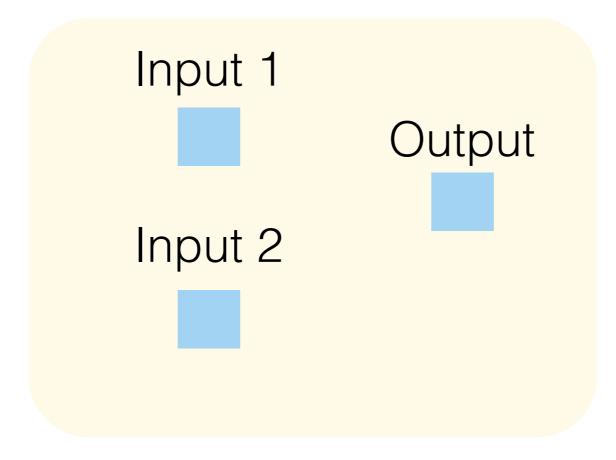
5,6

1,4



Pass 2





Pass 2

Output

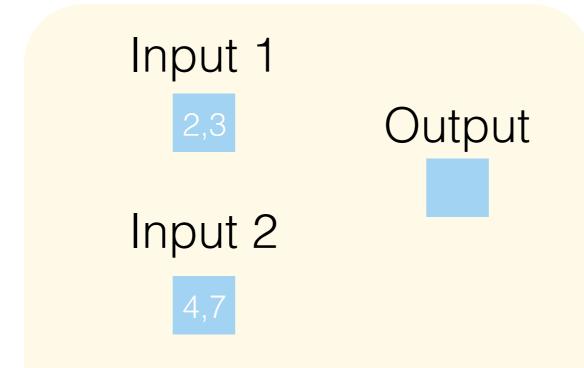
4,6

8,9

5,5

6,6

1,2



Pass 2

Output

4,6

8,9

5,5

6,6

1,2

4,4

Input 1



Output

2,3

Input 2

Pass 2

Output

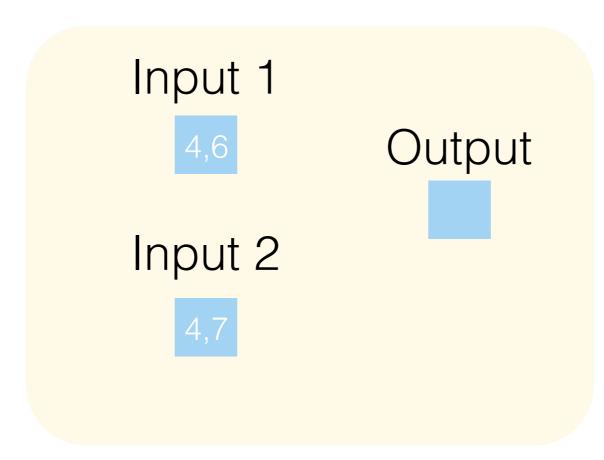
2,3

8,9

5,5

6,6

1,2



Pass 2

Output

2,3

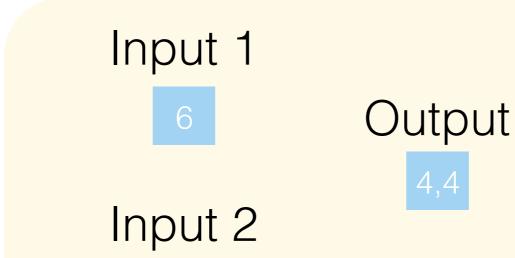
```
8,9
```

5,5

6,6

1,2

4,4



7

Pass 2

Output

2,3

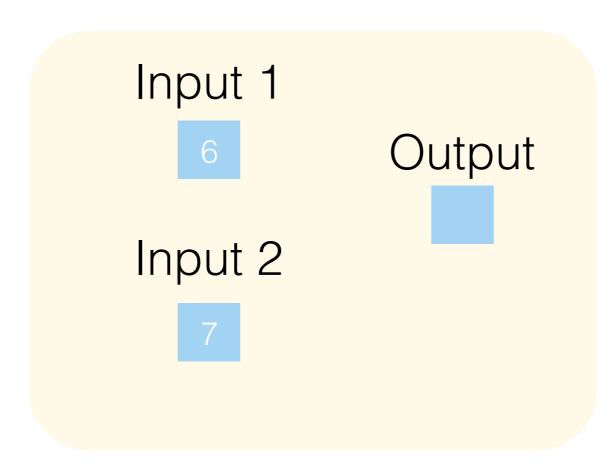
4,4



5,5

6,6

1,2



Pass 2

Output

2,3

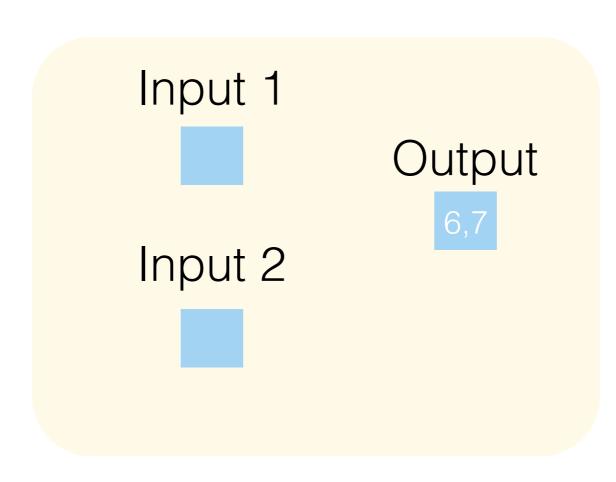
4,4



5,5

6,6

1,2



Pass 2

Output

2,3

4,4

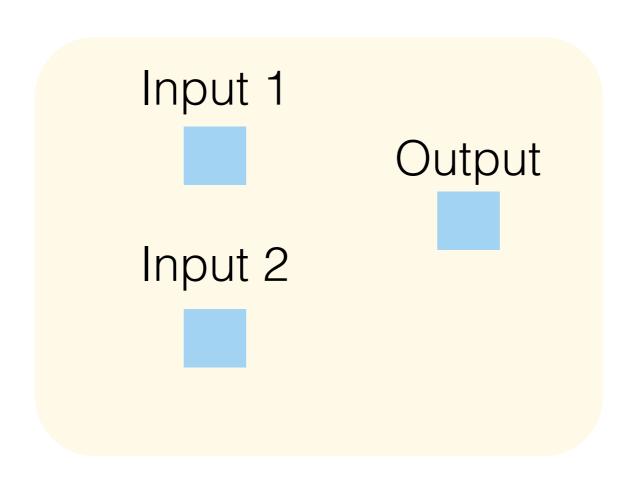
6,7

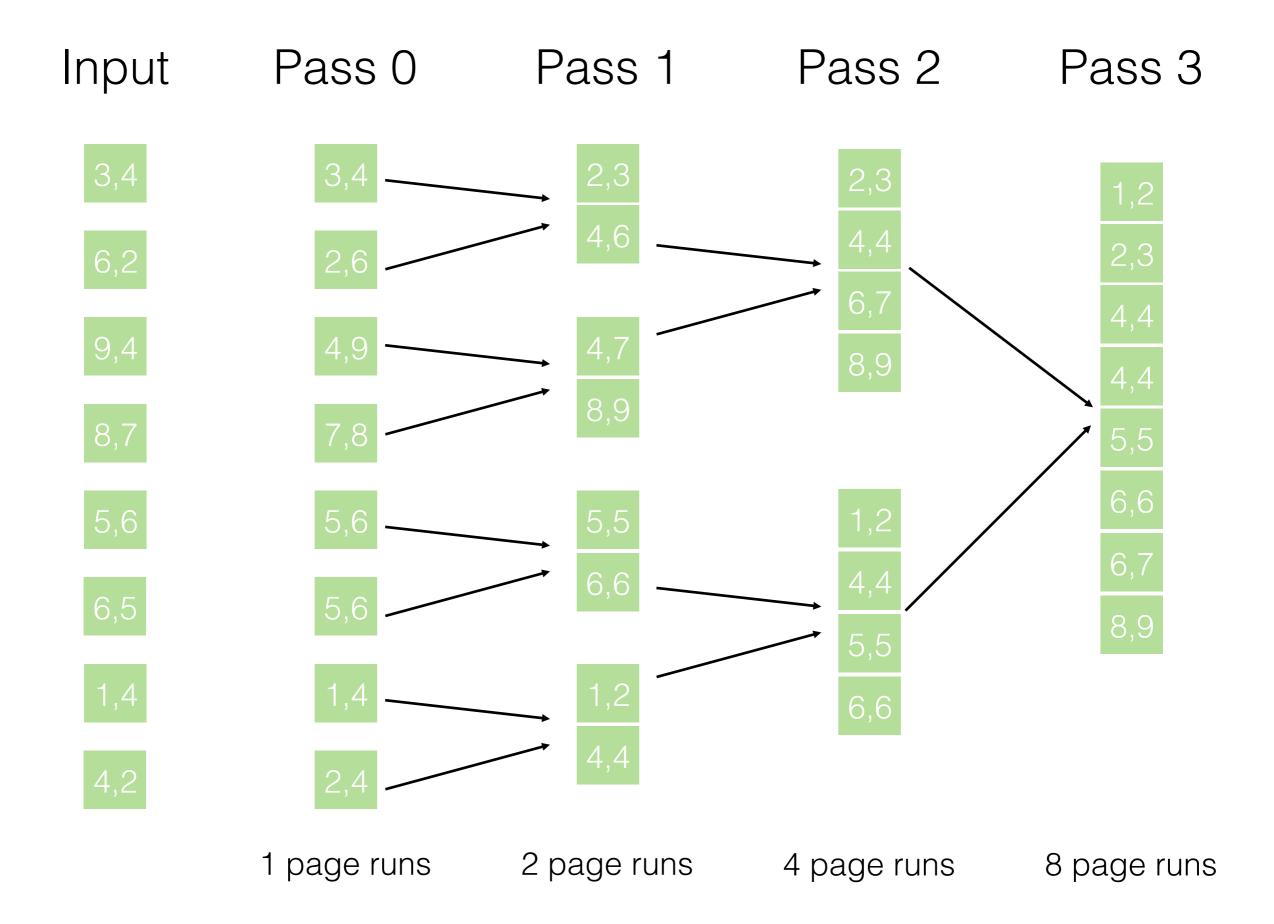
8,9

5,5

6,6

1,2





2-Way Merge Sort

Number of Passes:

```
• \lceil \log_2 N \rceil + 1
```

Number of I/O's:

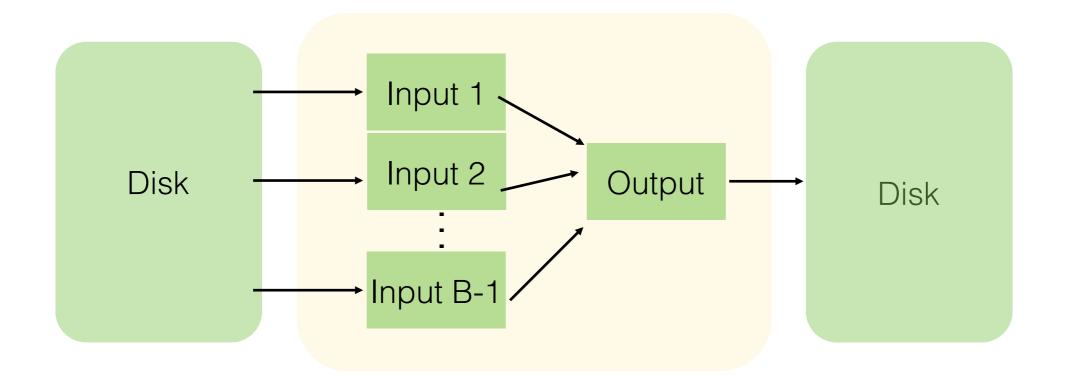
```
• 2N * (\lceil \log_2 N \rceil + 1)
```

Generalized Merge Sort

- Pass 0 (Conquer):
 - Sort B page at a time in memory
- Pass 1, 2, ..., etc. (Merge):
 - Merge B-1 runs

Generalized Merge Sort

Credits to Michelle Nguyen



Buffer size of B pages

Pass 0

Output

3,4

6,2

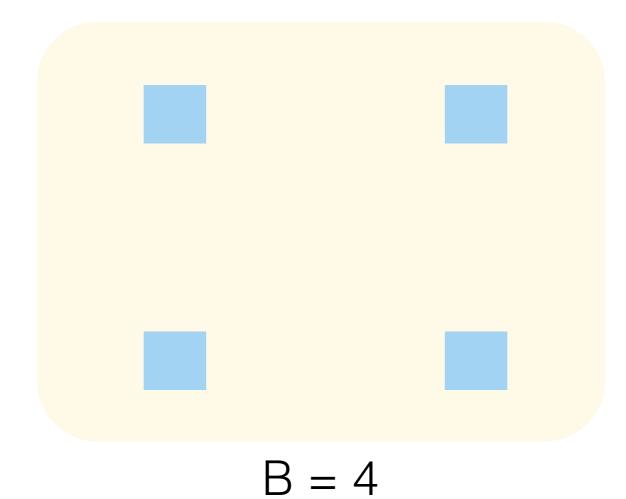
1.2

8,7

5

6,5

1,4



Pass 0

Output

5,3

1,2

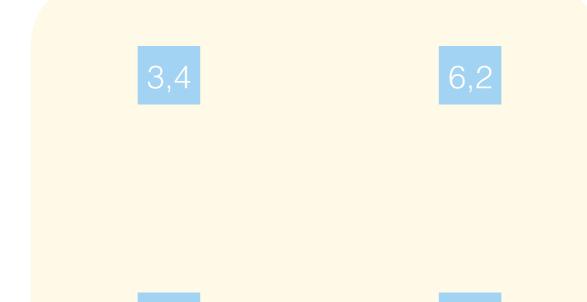
3,3

9,2

5,6

6,5

1,4



$$B = 4$$

Pass 0

Output

(Use a sort algorithm from 61B!)



1,2

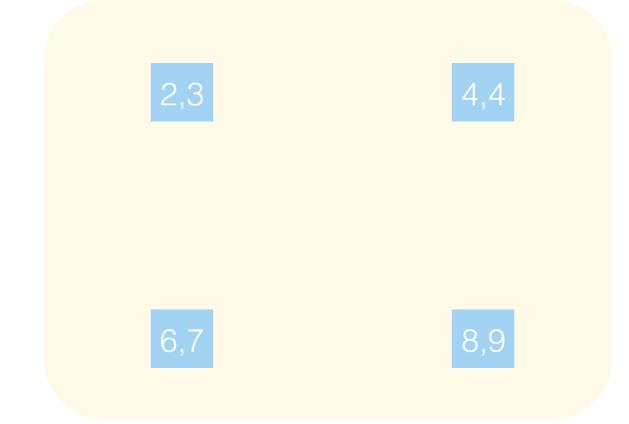
3,3

9,2

5,6

6,5

1,4



$$B = 4$$

Pass 0

Output

2,3

4,4

6,7

8,9

5.2

1,2

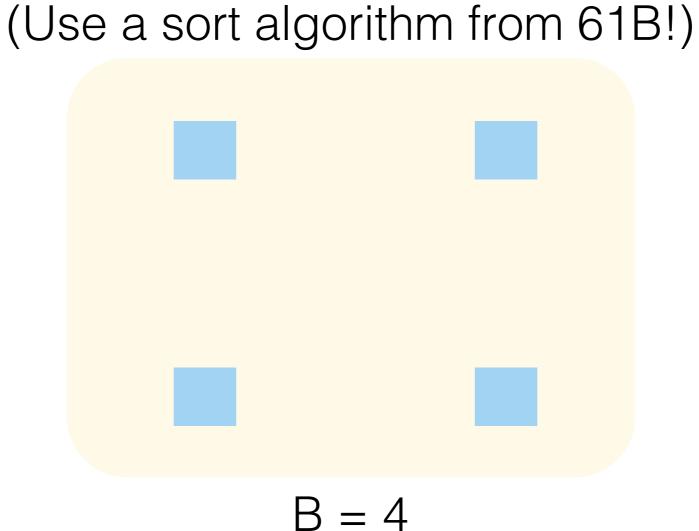
3,3

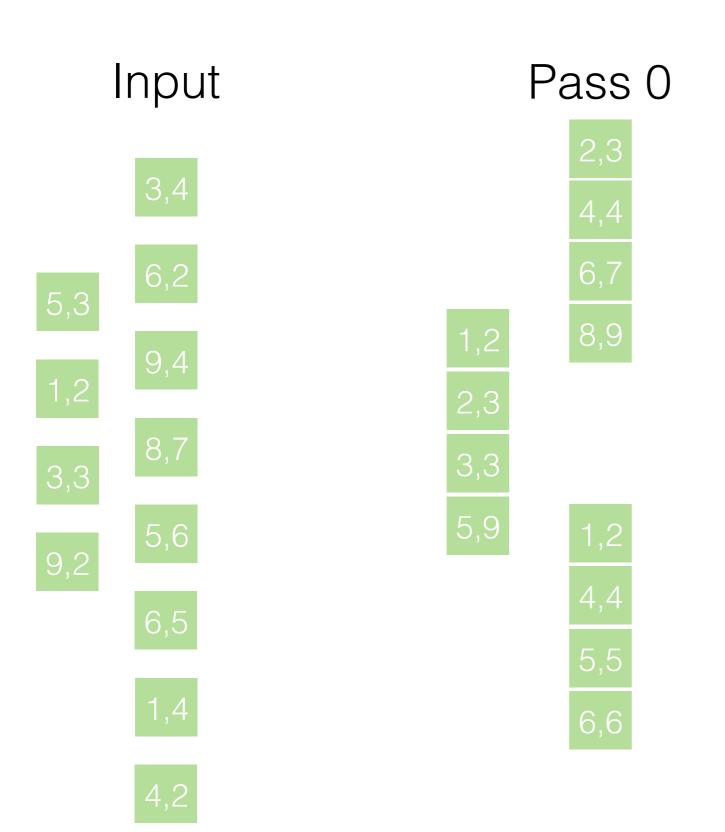
9,2

5,6

6,5

1,4





4 page runs

Pass 1

Output



1,2

2,3

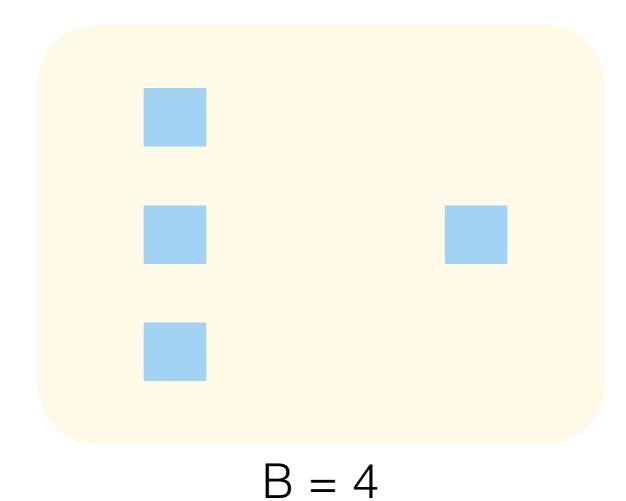
3,3

5,9

1 0

4,4

5,5



Pass 1

Output

4,4 6,7 8,9 3,3 5,9 4,4

```
a = 4
```

Pass 1

Output

```
\begin{bmatrix} 2, 3 \\ 2 \\ \end{bmatrix}
\begin{bmatrix} 1, 1 \\ 2 \end{bmatrix}
\begin{bmatrix} 3 \\ 4 \end{bmatrix}
```

Pass 1

Output

1,1

```
4,4
6.7
```

8.9

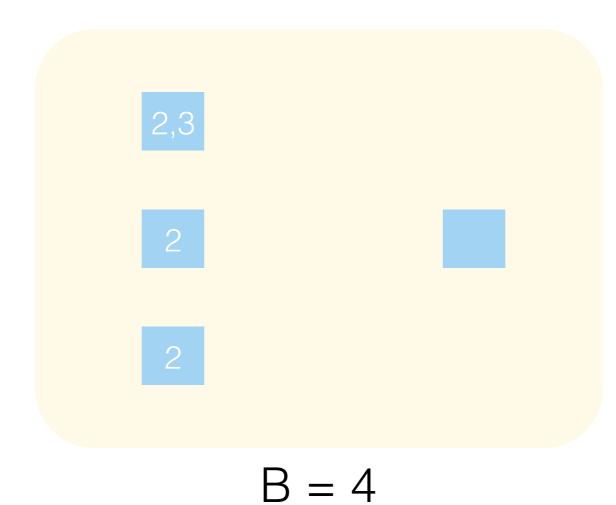
2,3

3,3

5,9

4,4

5,5



Pass 1

Output

1,1

```
6,7
```

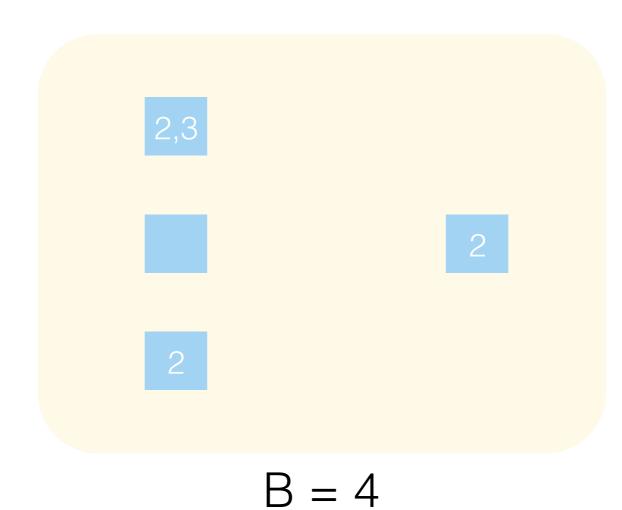
8,

3.3

5,9

4,4

5,5



Pass 1

Output

1,1

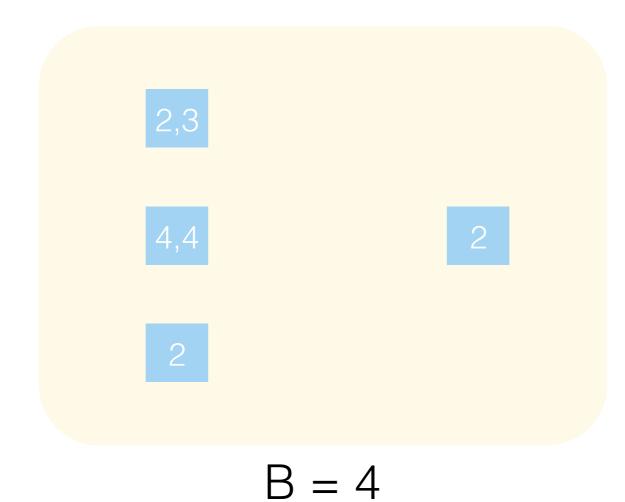
```
4,4
6.7
```

89

2,3

3,3





Pass 1

Output

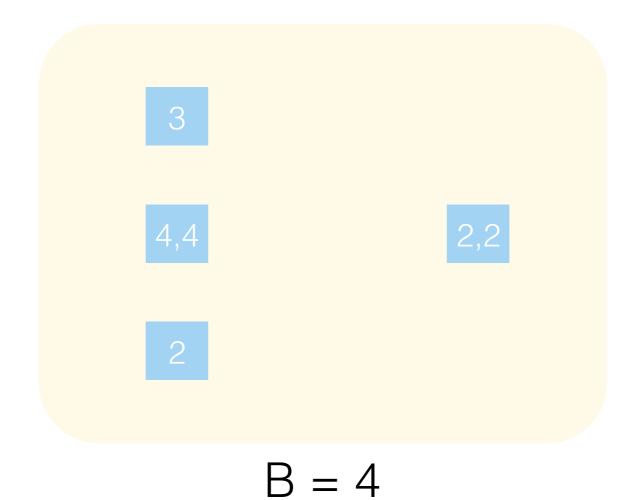
1,1

```
6,7
```

2,3

3.3



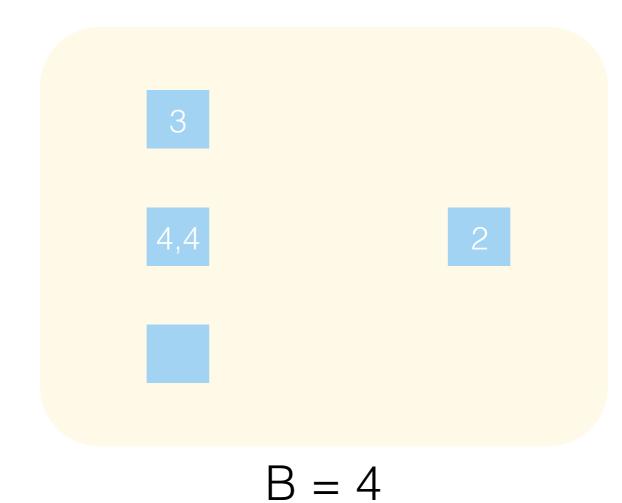


Pass 1

Output







Pass 1

Output

1,

2,2



8,9

3,3



$$B = 4$$

Generalized Merge Sort

Number of Passes:

```
• \lceil \log_{B-1} \lceil N/B \rceil \rceil + 1
```

Number of I/O's:

```
• 2N * (\lceil \log_{B-1} \lceil N/B \rceil \rceil + 1)
```

Worksheet #6, 7

6. List the differences between 2-way external merge sort and general external merge sort:

- 6. List the differences between 2-way external merge sort and general external merge sort:
 - Sorting: 2-way only utilizes 2 input buffers, general utilizes B-1
 - During pass 0, 2-way only uses 1 page to sort files.
 notice how applying B=3 to the general equation
 doesn't yield the correct # of passes. The general
 external merge sort uses all B pages in its buffer to sort
 the initial runs in pass 0.

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

1. How many pages can your buffer hold?

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

1. How many pages can your buffer hold?

640 KB * (1 page / 64 KB) = 10 pages

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

2. How many pages are in a 4 MB file?

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

2. How many pages are in a 4 MB file?

4 MB * (1024 KB / 1MB) * (1 page / 64 KB) = 64 pages

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

3. How many passes would it take to externally merge sort a 4 MB file?

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

3. How many passes would it take to externally merge sort a 4 MB file?

```
\lceil \log_{B-1} \lceil N/B \rceil \rceil + 1
= \lceil \log_{10-1} \lceil 64/10 \rceil \rceil + 1
= \lceil \log_{9} 7 \rceil + 1
= 1 + 1
= 2 passes
```

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

4. How many I/O's are needed to to externally merge sort a 4 MB file?

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

4. How many I/O's are needed to to externally merge sort a 4 MB file?

```
(# of passes) * 2 * (# of pages in file)
= 2 * 2 * 64
= 256 I/O's
```

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

5. What is the maximum file size that can be sorted with just 2 passes in this system?

External Sorting Exercises

7. Your system has 640 KB of memory allocated for the buffer for external sorting and you have infinite space for scratch disks. Each page holds 64 KB of data.

Note: 1024 KB = 1 MB.

5. What is the maximum file size that can be sorted with just 2 passes in this system?

```
(# of buffer pages) ( # of buffer pages - 1 )
= 10 * 9
= 90 pages
= 5760 KB ~= 5.6 MB
```

External Hashing

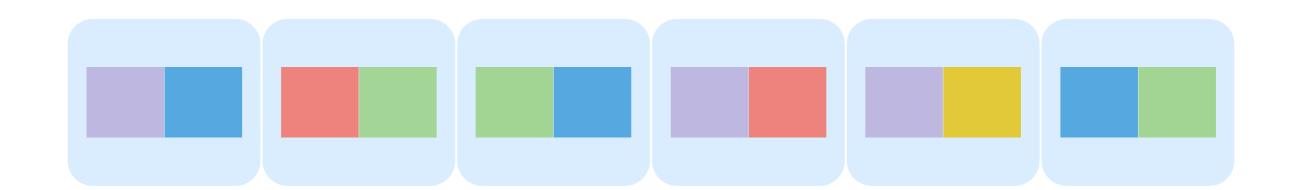
External Hashing

Want to aggregate data that does not fit in memory

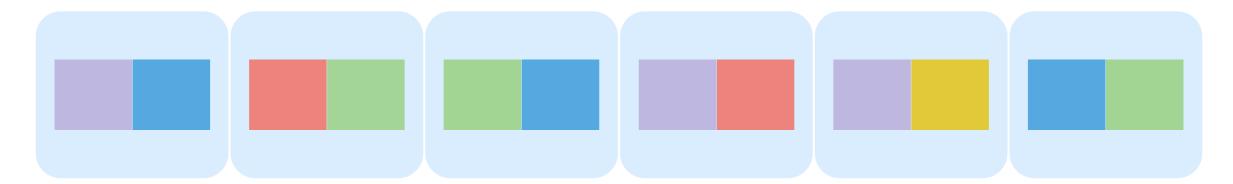
Aggregating Colors

Credits to Michelle Nguyen

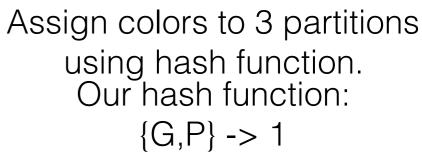
- Goal: Group squares by color
- Setup: 12 squares, 2 can fit per page. We can hold 8 squares in memory.
- N=6, B=4

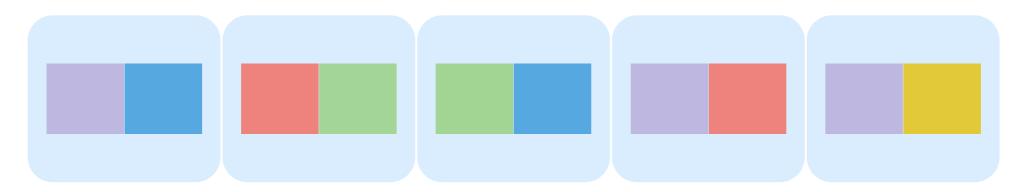


- Read all pages in, hash to B-1 partitions/buckets so that each group guaranteed to be in same partition.
- May not be a whole partition for each group.
- # I/O's = 2N

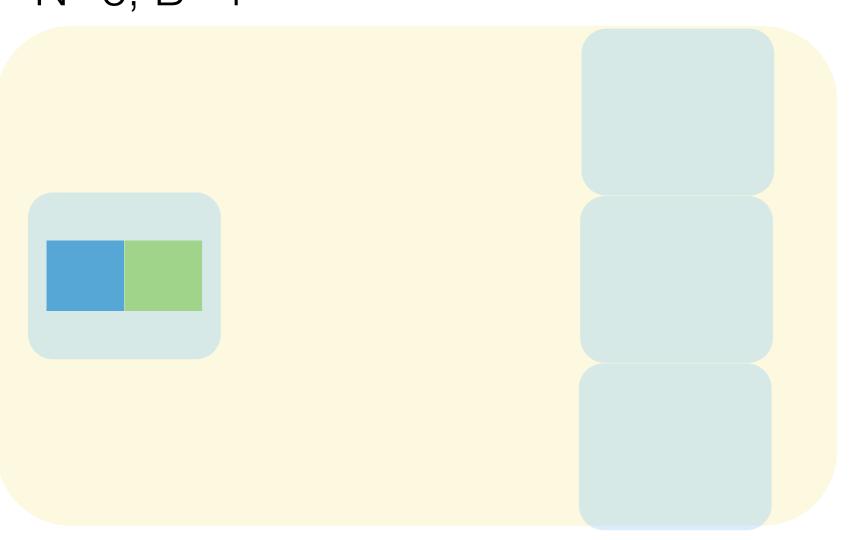


$$N=6, B=4$$



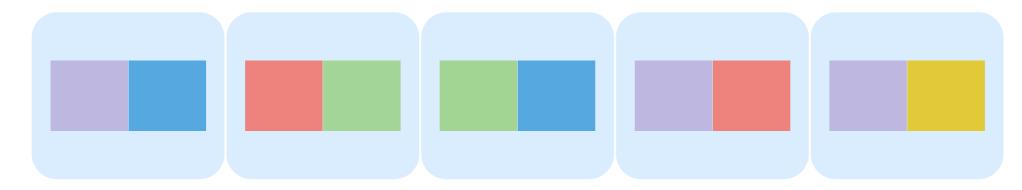


$$N=6, B=4$$

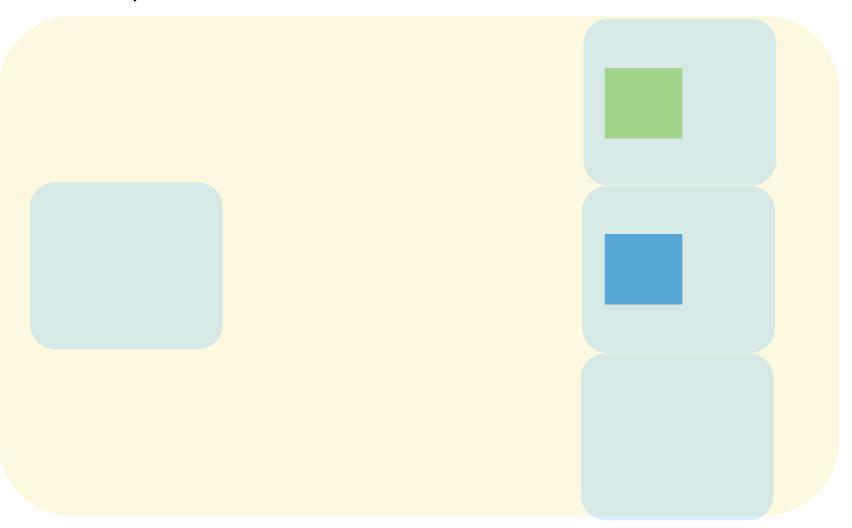


$${G,P} \rightarrow 1$$

 ${B} \rightarrow 2$
 ${R, Y} \rightarrow 3$

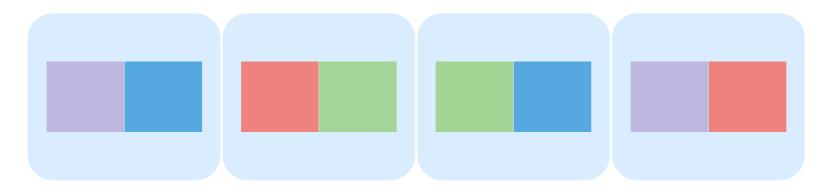


$$N=6, B=4$$

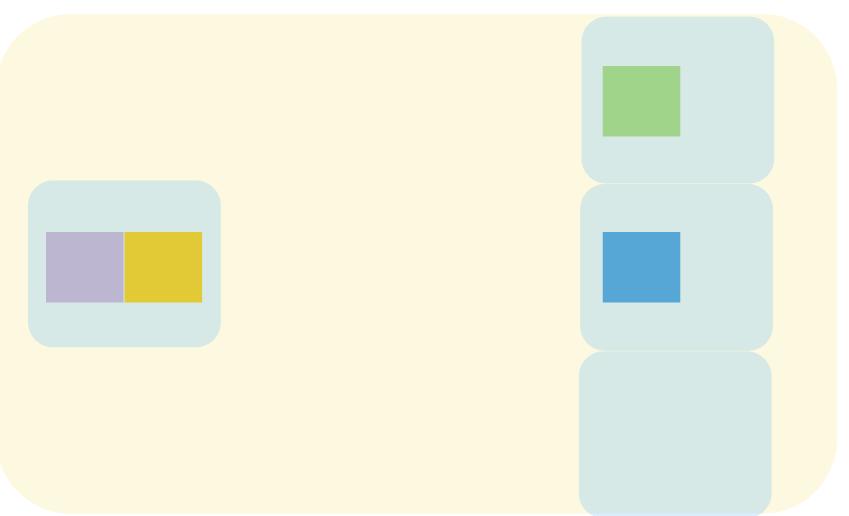


$${G,P} \rightarrow 1$$

 ${B} \rightarrow 2$
 ${R, Y} \rightarrow 3$

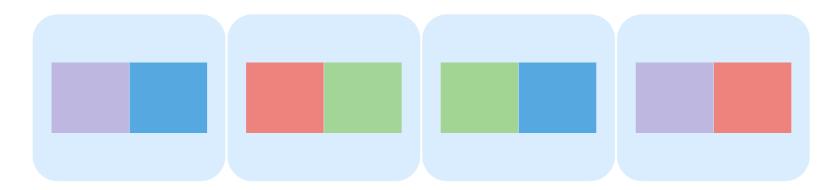


$$N=6, B=4$$

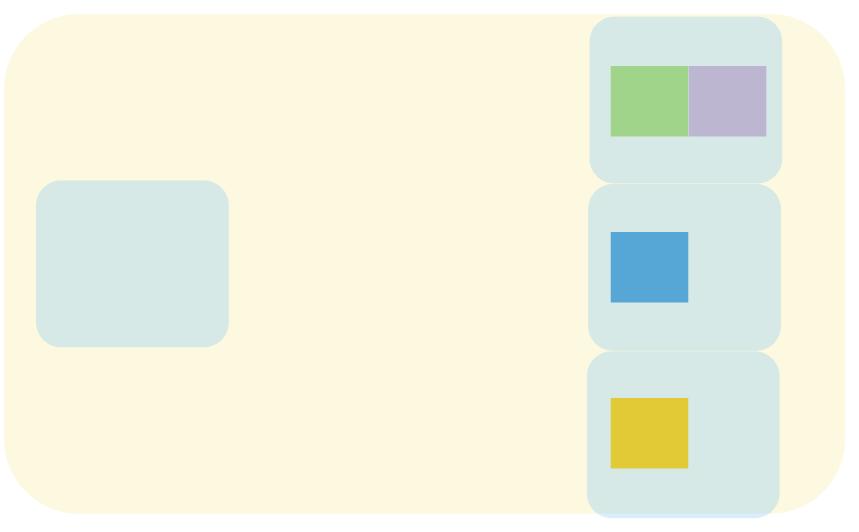


$${G,P} \rightarrow 1$$

 ${B} \rightarrow 2$
 ${R, Y} \rightarrow 3$

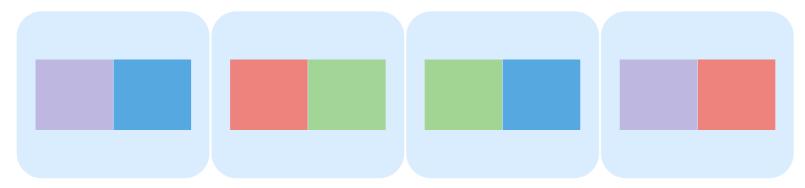


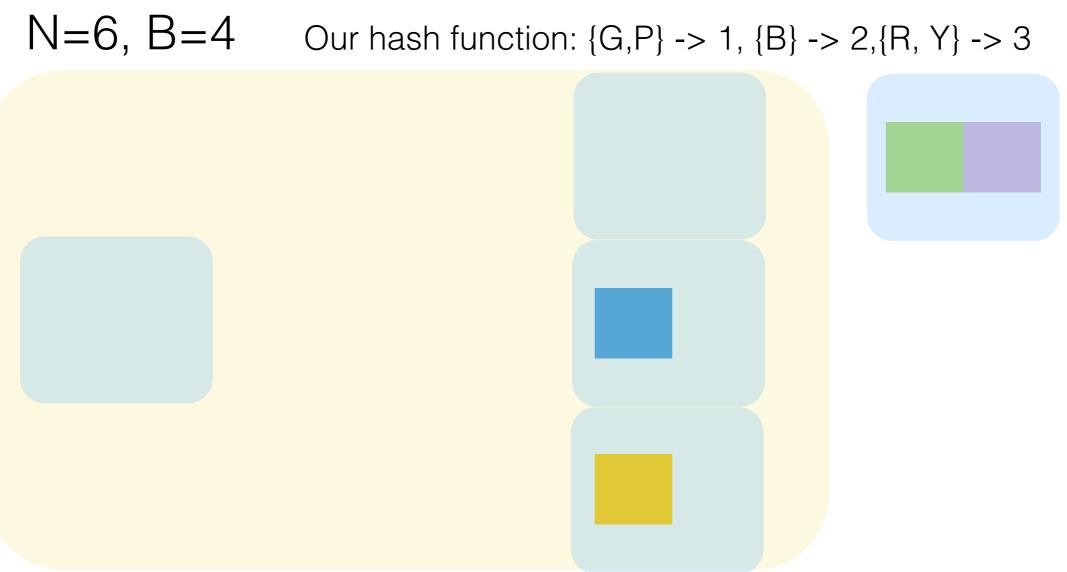
$$N=6, B=4$$

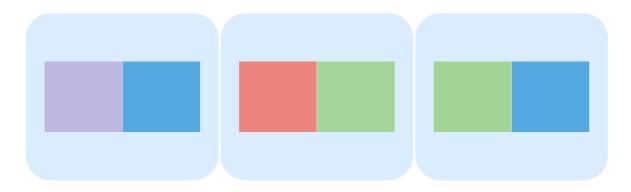


$${G,P} \rightarrow 1$$

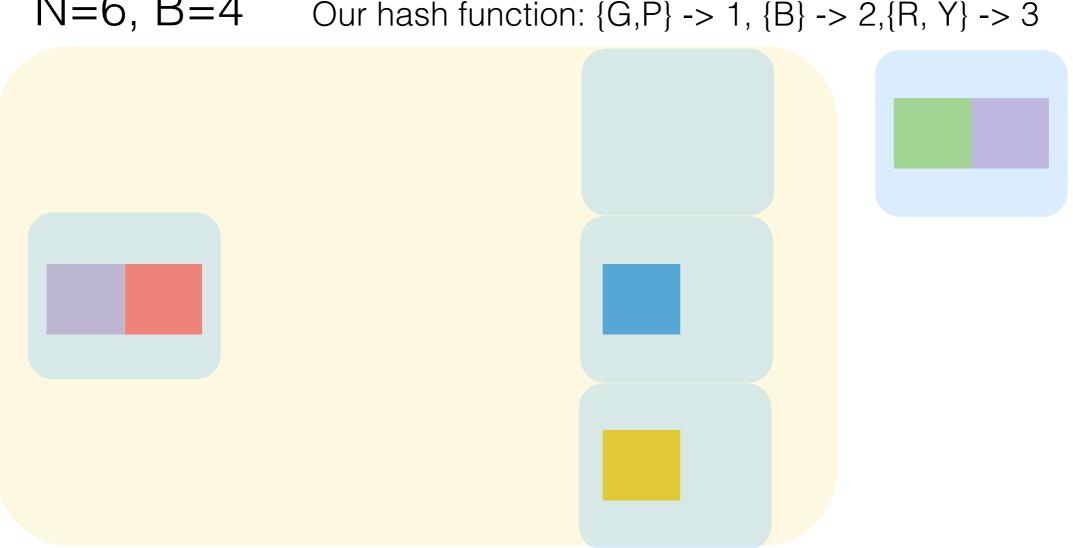
 ${B} \rightarrow 2$
 ${R, Y} \rightarrow 3$

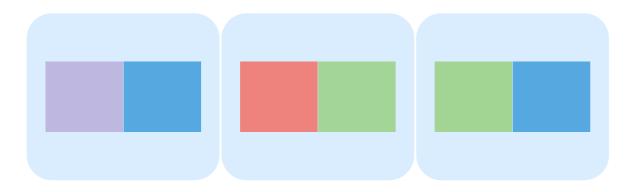




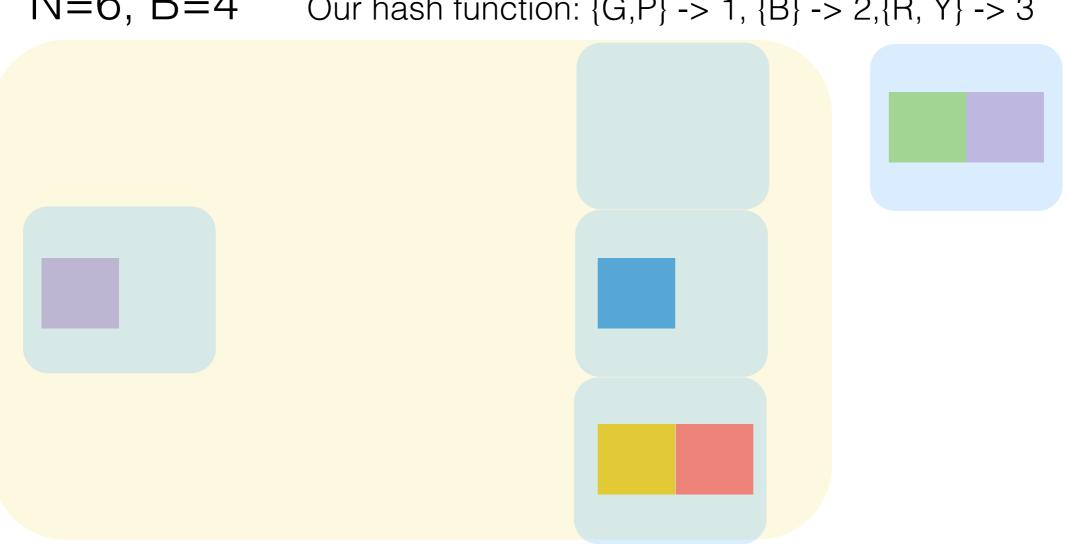


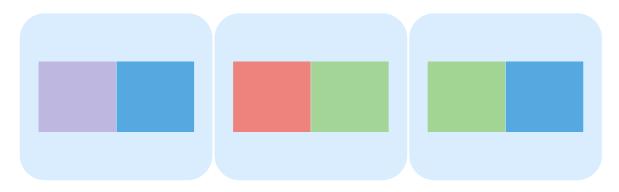
N=6, B=4Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$

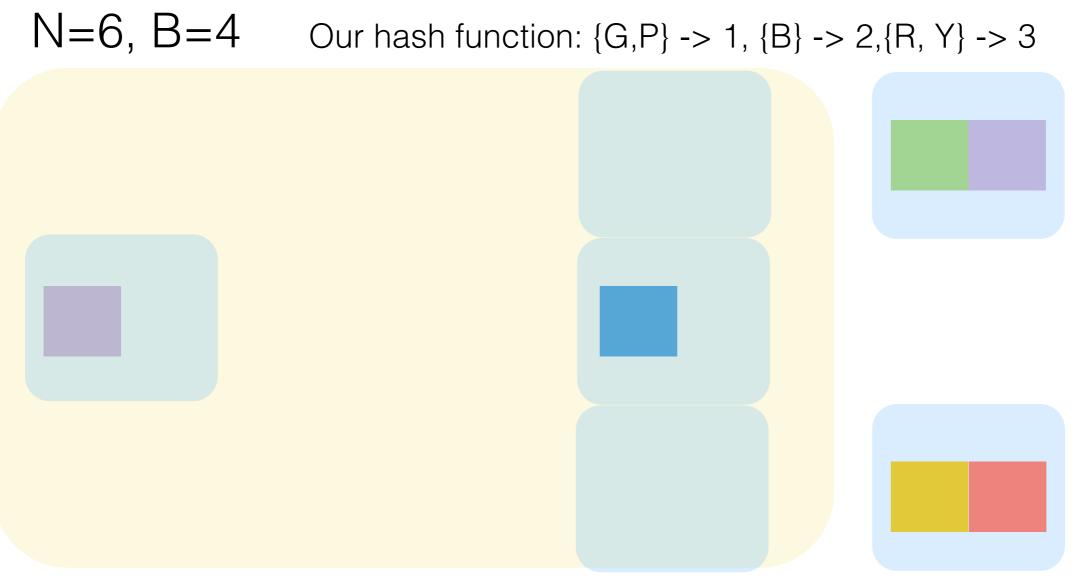


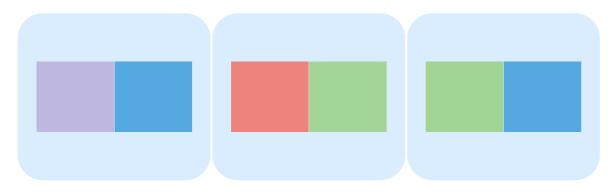


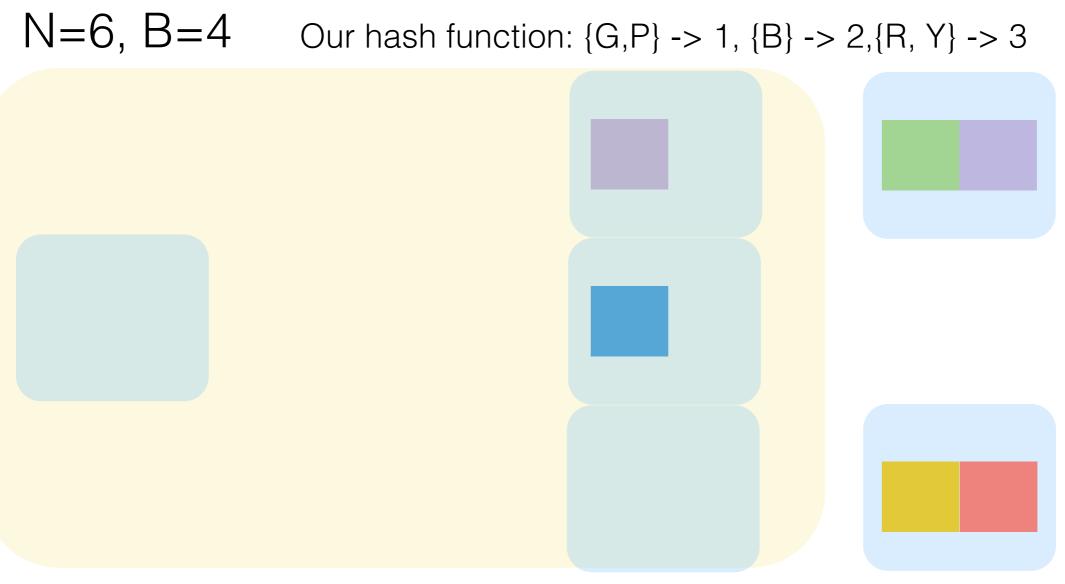
N=6, B=4Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$

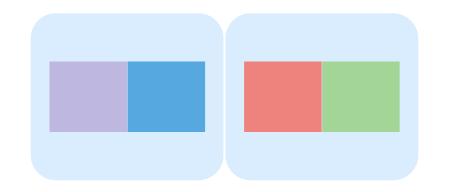




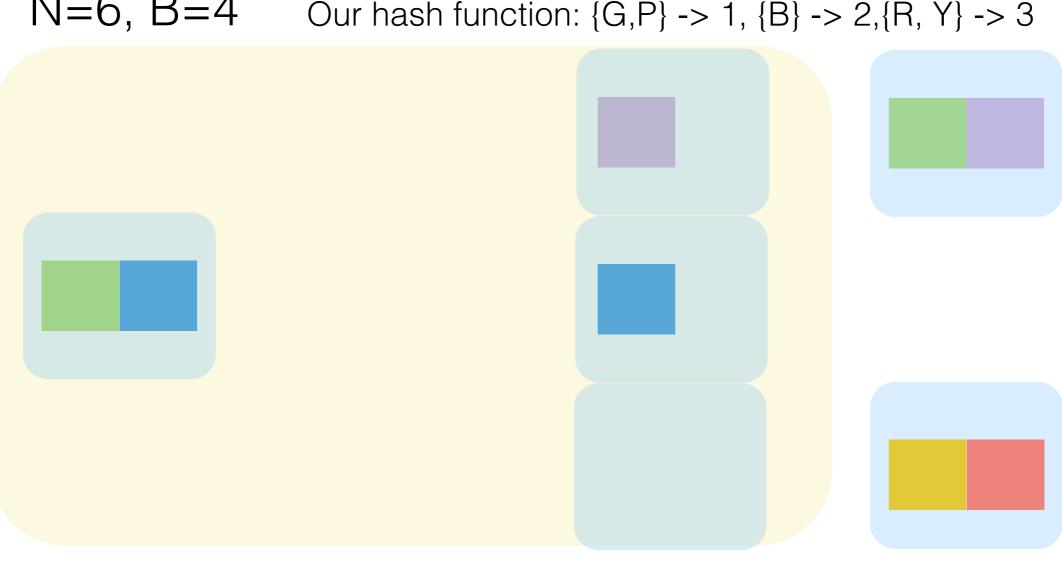


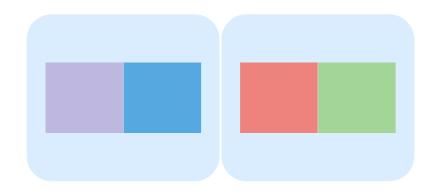




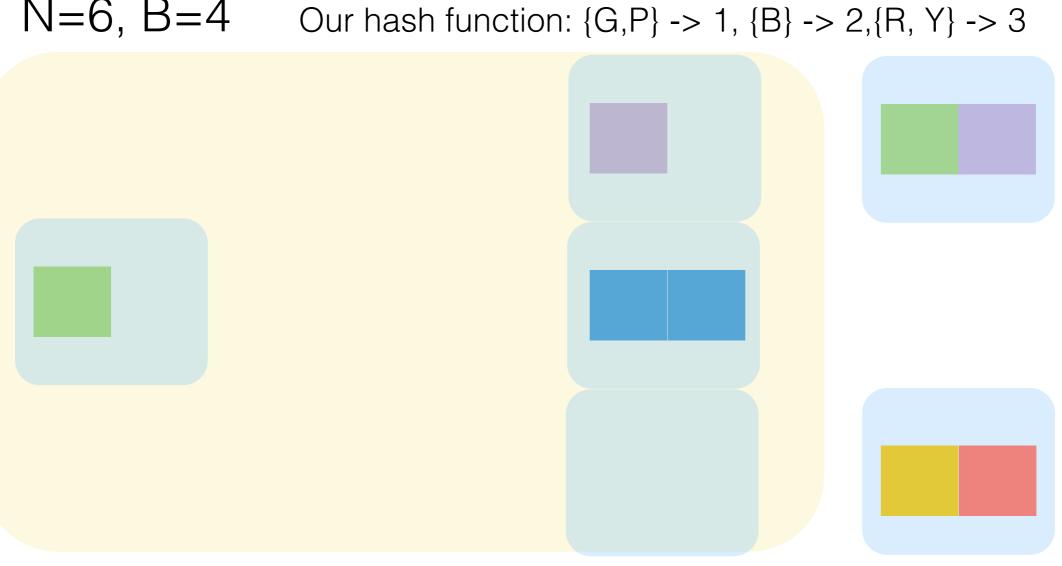


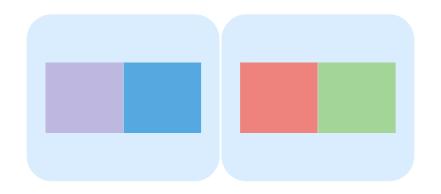
N=6, B=4Our hash function: {G,P} -> 1, {B} -> 2,{R, Y} -> 3

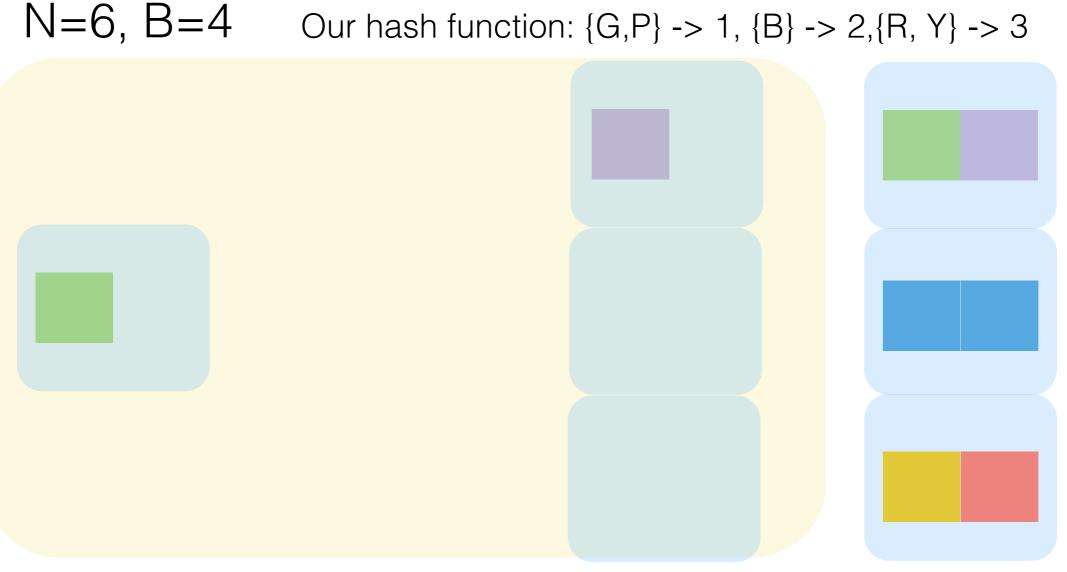


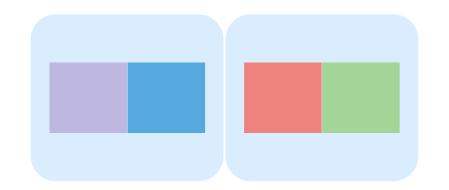


N=6, B=4

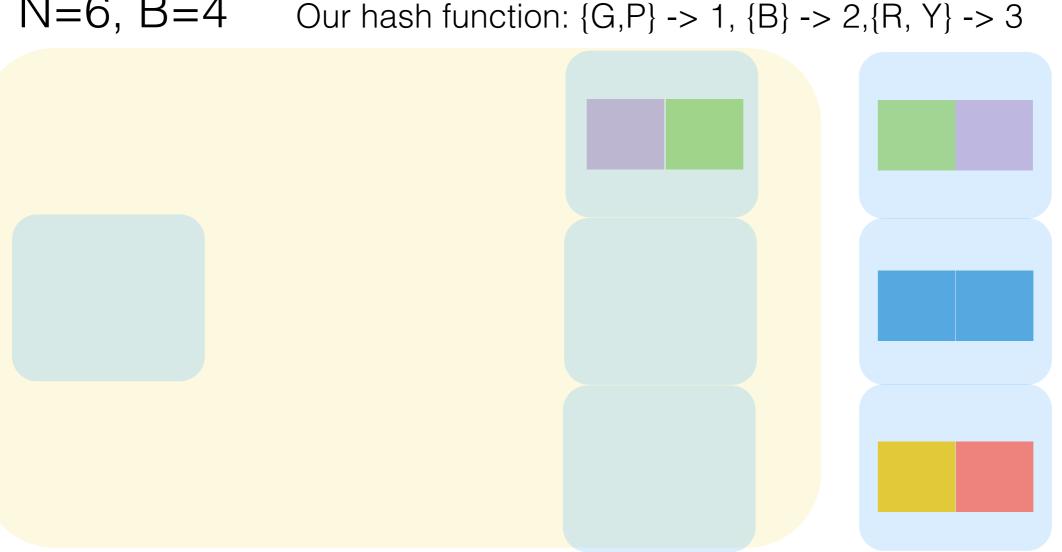


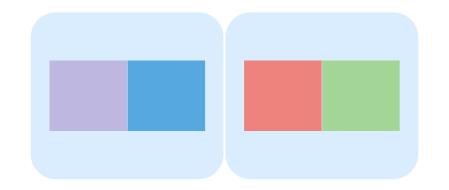




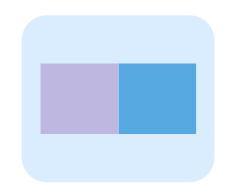


N=6, B=4Our hash function: {G,P} -> 1, {B} -> 2,{R, Y} -> 3



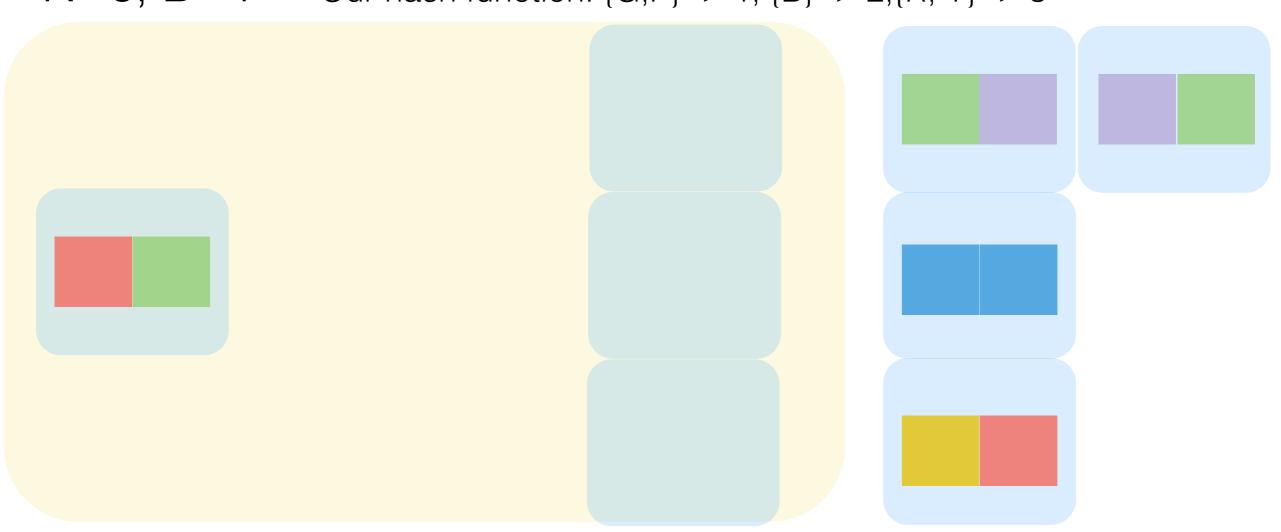


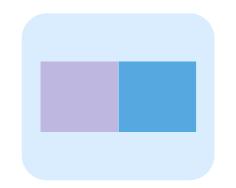
N=6, B=4 Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$



N=6, B=4

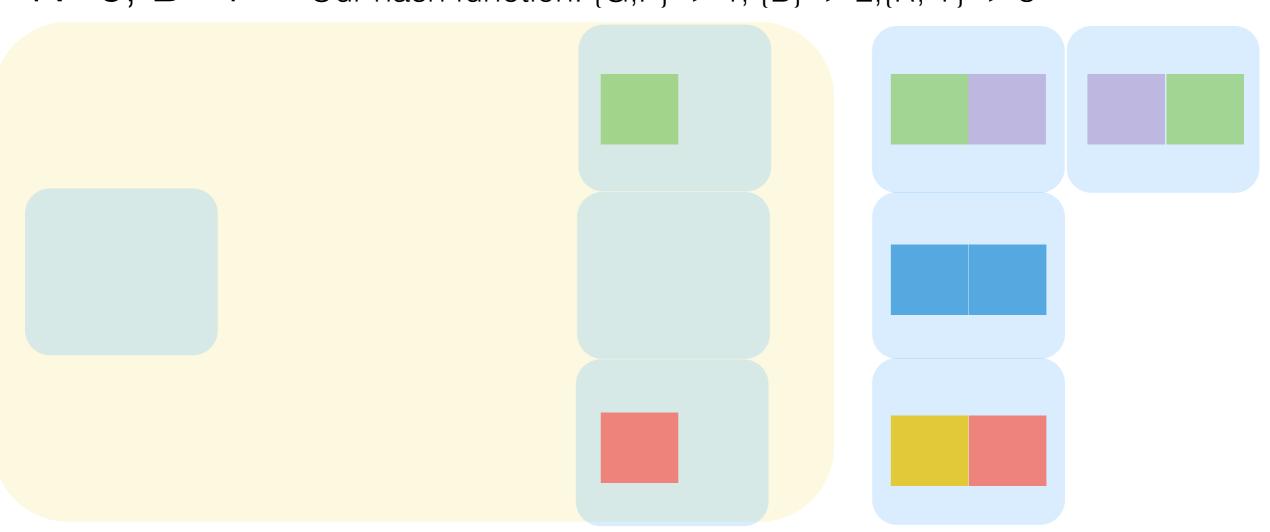
Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$



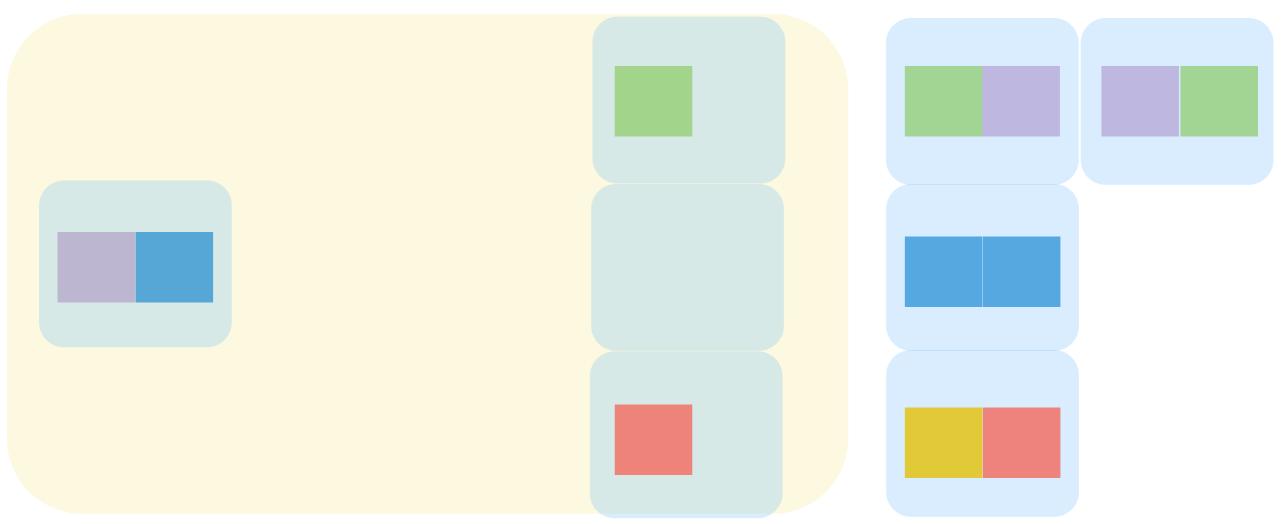


N=6, B=4

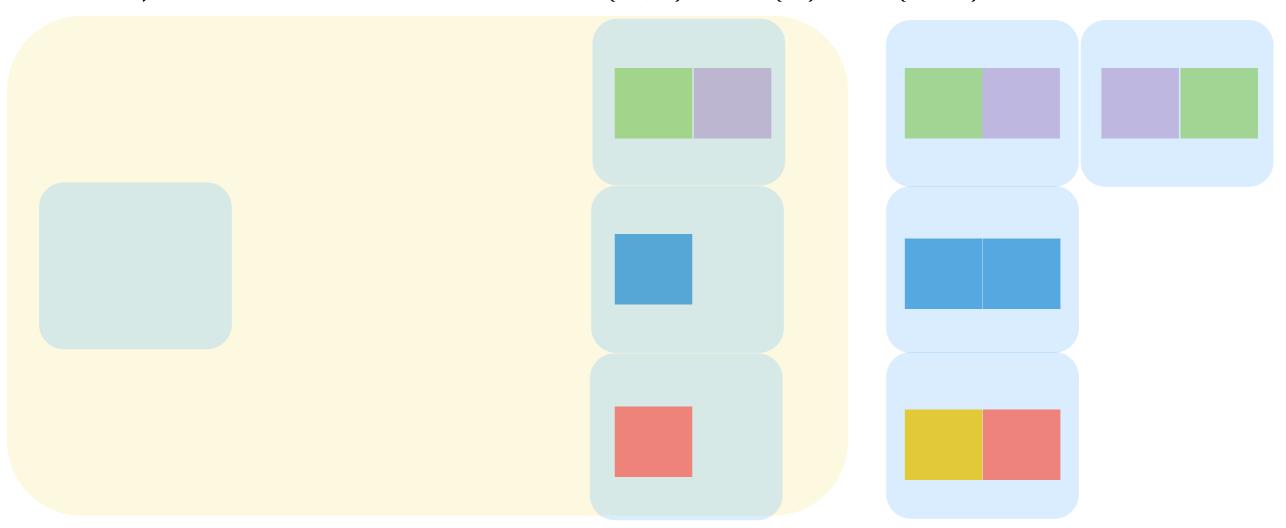
Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$



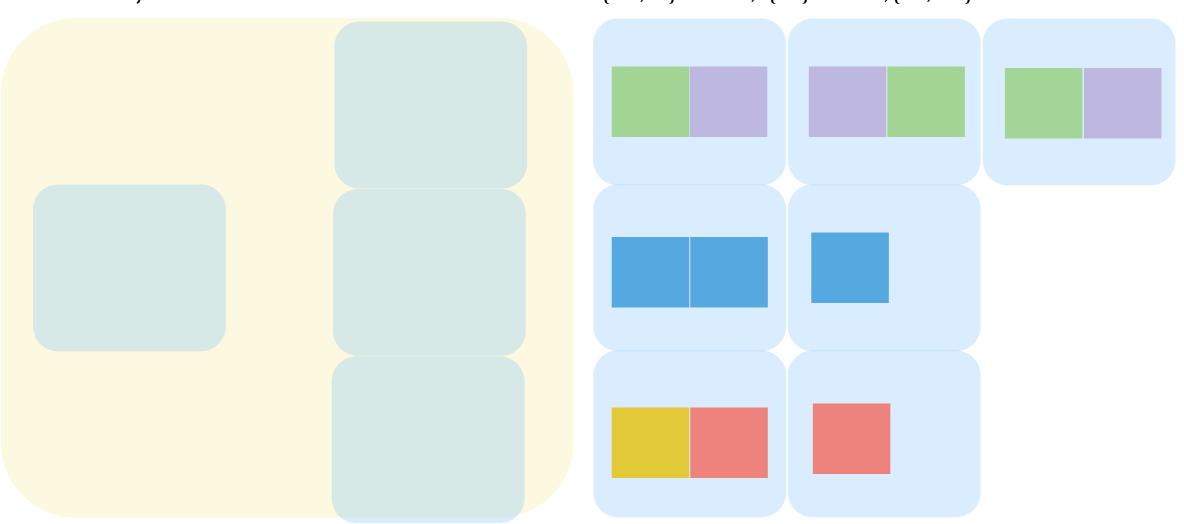
N=6, B=4 Our hash function: $\{G,P\} -> 1$, $\{B\} -> 2$, $\{R,Y\} -> 3$



N=6, B=4 Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$



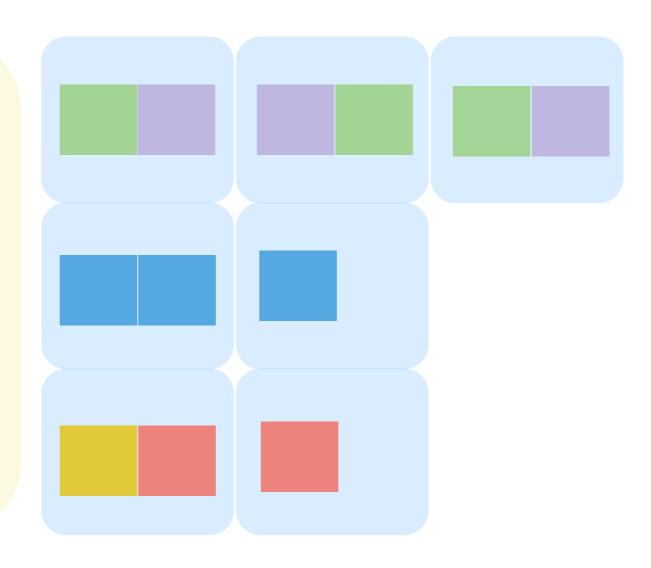
N=6, B=4 Our hash function: $\{G,P\} \rightarrow 1$, $\{B\} \rightarrow 2$, $\{R,Y\} \rightarrow 3$



- Rehash each partition.
- For a partition to fit in memory, it can only have B pages.
- To hash larger tables, use the partition algorithm recursively until the partition fits into memory
- # I/O's = 2N

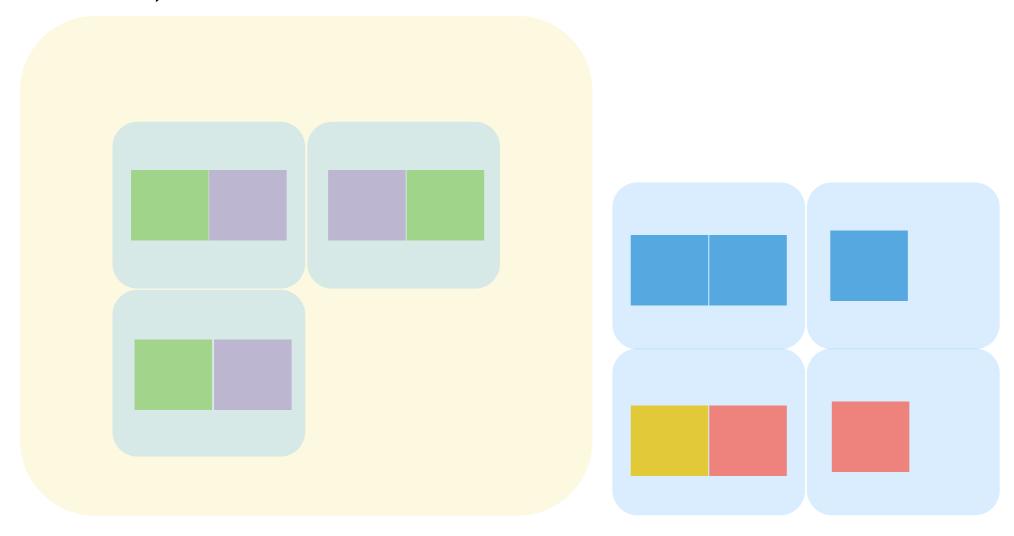
Create in-memory table for each partition.

$$N=6, B=4$$



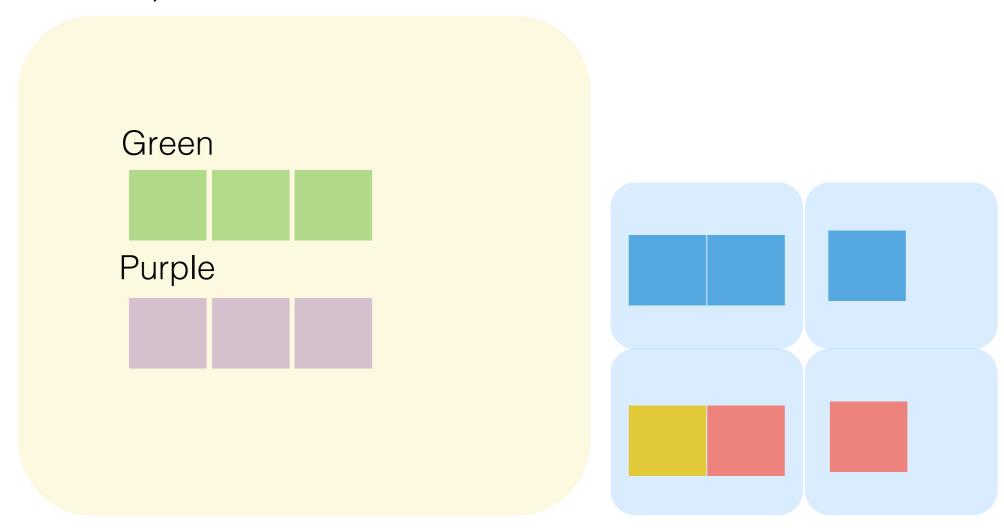
Create in-memory table for each partition.

$$N=6, B=4$$



Create in-memory table for each partition.

$$N=6, B=4$$



Worksheet #3, 4, 5

3. Why can we process B * (B - 1) pages of data with external hashing in just two passes (divide and conquer phases)?

3. Why can we process B * (B - 1) pages of data with external hashing in just two passes (divide and conquer phases)?

Our main limitation is how big the partitions can be after the partition hashing. Since we need to be able to read in the whole partition into memory, each partition can be at most B pages big.

4. If you're processing exactly B * (B - 1) pages of data, is it likely that you'll have to perform recursive external hashing? Why?

4. If you're processing exactly B * (B - 1) pages of data, is it likely that you'll have to perform recursive external hashing? Why?

You would have to have an absolutely perfect hash function that evenly distributes any record into the B-1 partitions. This is almost impossible in practice. Rather, we should expect that some partitions may be larger than B after partition hashing.

5. While you recursively perform external hashing, you reuse the same hash functions for partitioning. What's the problem with this?

5. While you recursively perform external hashing, you reuse the same hash functions for partitioning. What's the problem with this?

The partition that is too big to fit in memory will still be too big to fit in memory if we maintain the same partition hashing strategy.